

## ▼ Install easy-vqa library

```
# this package has all the data needed to train our model
!pip install easy-vqa

Collecting easy-vqa
  Downloading easy_vqa-1.0-py3-none-any.whl (3.2 MB)
     |-----| 3.2/3.2 MB 2.2 MB/s eta 0:00:00
Installing collected packages: easy-vqa
Successfully installed easy-vqa-1.0
```

## Inputs to model

We would need the following inputs to our model in the form of

- [Image, Question, Answer]

For train and test datasets these will be the variable inputs

- [train\_X\_ims, train\_X\_seqs, train\_Y]- train image-question-answer input
- [test\_X\_ims, test\_X\_seqs, test\_Y]- test image-question-answer input

We would also need a

- embedding\_size - size of embedding for input node count
- im\_shape - shape of image for input node count
- num\_answers - number of answers for output layer node count

## ▼ Loading and preprocessing Images

```
import numpy as np
import torch
import torchvision
import torchvision.transforms as T
from torchvision.transforms import functional as F
from PIL import Image
from easy_vqa import get_train_questions, get_test_questions, get_train_image_paths, get_test_image_paths, get_answers

def load_and_process_image(image_path):
    # Loads image from path and converts to Tensor, you can also reshape the im
    im = Image.open(image_path)
    im = F.to_tensor(im)
    return im

def read_images(paths):
    # paths is a dict mapping image ID to image path
    # Returns a dict mapping image ID to the processed image
    ims = {}
    for image_id, image_path in paths.items():
        ims[image_id] = load_and_process_image(image_path)
    return ims

print('--- Reading/processing images from image paths of the vqa library ---\n')
train_ims = read_images(get_train_image_paths())
test_ims = read_images(get_test_image_paths())
im_shape = train_ims[0].shape
print(f'Read {len(train_ims)} training images and {len(test_ims)} testing images.')
print(f'Each image has shape {im_shape}.')

print('\n--- Creating model input images...')
train_qs, train_answers, train_image_ids = get_train_questions()
test_qs, test_answers, test_image_ids = get_test_questions()
train_X_ims = np.array([train_ims[id] for id in train_image_ids])
test_X_ims = np.array([test_ims[id] for id in test_image_ids])

--- Reading/processing images from image paths of the vqa library ---

Read 4000 training images and 1000 testing images.
Each image has shape torch.Size([3, 64, 64]).
```

```

--- Creating model input images...
<ipython-input-2-c6865f93d0ef>:33: FutureWarning: The input object of type 'Tensor' is an array-like implementing one of the
  train_X_ims = np.array([train_ims[id] for id in train_image_ids])
<ipython-input-2-c6865f93d0ef>:33: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a 1
  train_X_ims = np.array([train_ims[id] for id in train_image_ids])
<ipython-input-2-c6865f93d0ef>:34: FutureWarning: The input object of type 'Tensor' is an array-like implementing one of the
  test_X_ims = np.array([test_ims[id] for id in test_image_ids])
<ipython-input-2-c6865f93d0ef>:34: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a 1
  test_X_ims = np.array([test_ims[id] for id in test_image_ids])

```

## ▼ Loading Questions and Answers

```

print('\n--- Reading questions...')
train_qs, train_answers, train_image_ids = get_train_questions()
test_qs, test_answers, test_image_ids = get_test_questions()
print(f'Read {len(train_qs)} training questions and {len(test_qs)} testing questions.')

```

```

print('\n--- Reading answers...')
all_answers = get_answers()
num_answers = len(all_answers)
print(f'Found {num_answers} total answers:')
print(all_answers)

```

```

--- Reading questions...
Read 38575 training questions and 9673 testing questions.

```

```

--- Reading answers...
Found 13 total answers:
['circle', 'green', 'red', 'gray', 'yes', 'teal', 'black', 'rectangle', 'yellow', 'triangle', 'brown', 'blue', 'no']



```

## ▼ Quick look at the dataset

```

import pandas as pd
df = pd.DataFrame(list(zip(train_qs, train_answers, train_image_ids)), columns=['Question', 'Answer', 'Image ID'])
df.head(10)

```

	Question	Answer	Image ID	
0	what is the blue shape?	rectangle	0	
1	what color is the shape?	blue	0	
2	does the image contain a rectangle?	yes	0	
3	is there a triangle in the image?	no	0	
4	is there a black shape?	no	0	
5	does the image not contain a gray shape?	yes	0	
6	is there a red shape in the image?	no	0	
7	does the image not contain a red shape?	yes	0	
8	is there not a blue shape?	no	0	
9	is there not a blue shape in the image?	no	0	

## ▼ Quick look at images

```

import torchvision.utils as utils
from torchvision import transforms

# print multiple images
# images = 1
# batch = torch.empty((images, 3, 64, 64))
# for i in range(images):
#     batch[i] = train_ims[i]

# Create a grid of images
id = 0
grid = utils.make_grid(train_ims[id], nrow=2)

```

```
# Convert the grid to a PIL image
image = transforms.ToPILImage()(grid)
# Show the image
image.show()
```

## ▼ Preprocessing Questions

```
! pip install -U sentence-transformers
from sentence_transformers import SentenceTransformer, util
st_model = SentenceTransformer('multi-qa-MiniLM-L6-cos-v1')

#Questions are encoded by calling model.encode()
train_X_seqs = st_model.encode(train_qs)
test_X_seqs = st_model.encode(test_qs)

# convert ndarray to tensor
train_X_seqs = torch.tensor(train_X_seqs, dtype=torch.float)
test_X_seqs = torch.tensor(test_X_seqs, dtype=torch.float)

print(f'\n\nThe shape of the binary vectors is : {train_X_seqs.shape}')
```

```
Collecting sentence-transformers
  Downloading sentence-transformers-2.2.2.tar.gz (85 kB)
    86.0/86.0 kB 3.0 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Collecting transformers<5.0.0,>=4.6.0 (from sentence-transformers)
  Downloading transformers-4.34.1-py3-none-any.whl (7.7 MB)
    7.7/7.7 MB 66.7 MB/s eta 0:00:00
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (4.66.1)
Requirement already satisfied: torch>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (2.1.0+cu118)
Requirement already satisfied: torchvision in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (0.16.0+cu118)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (1.23.5)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (1.2.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (1.11.3)
Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-packages (from sentence-transformers) (3.8.1)
Collecting sentencepiece (from sentence-transformers)
  Downloading sentencepiece-0.1.99-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.3 MB)
    1.3/1.3 MB 86.1 MB/s eta 0:00:00
Collecting huggingface-hub>=0.4.0 (from sentence-transformers)
  Downloading huggingface_hub-0.18.0-py3-none-any.whl (301 kB)
    302.0/302.0 kB 39.2 MB/s eta 0:00:00
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.4.0->sentence-transformers) (3.12.2)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.4.0->sentence-transformers) (2023.6.0)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.4.0->sentence-transformers) (2.31.0)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.4.0->sentence-transformers) (6.0.1)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.4.0->sentence-transformers) (4.5.0)
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.4.0->sentence-transformers) (23.1)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->sentence-transformers) (1.11.1)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->sentence-transformers) (3.1)

```

## ▼ Preprocessing Answers

```

Collecting transformers<5.0.0,>=4.6.0 (from sentence-transformers)
  Downloading transformers-4.34.1-py3-none-any.whl (7.7 MB)
    7.7/7.7 MB 66.7 MB/s eta 0:00:00
print('\n--- Creating model outputs...')

train_answer_indices = np.array([all_answers.index(a) for a in train_answers])
test_answer_indices = np.array([all_answers.index(a) for a in test_answers])

#creating a 2D array filled with 0's
train_Y = np.zeros((train_answer_indices.size, train_answer_indices.max()+1), dtype=int)
test_Y = np.zeros((test_answer_indices.size, test_answer_indices.max()+1), dtype=int)

#replacing 0 with a 1 at the index of the original array
train_Y[np.arange(train_answer_indices.size),train_answer_indices] = 1
test_Y[np.arange(test_answer_indices.size),test_answer_indices] = 1

# finally convert the label vectors to tensor and fix the data type so it wouldnt error in the fully connected lay
train_Y = torch.tensor(train_Y, dtype=torch.float)
test_Y = torch.tensor(test_Y, dtype=torch.float)

print(f'Example model output: {train_Y[0]}')
print(f'data type {type(train_Y)}')

--- Creating model outputs...
Example model output: tensor([0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0.])
data type <class 'torch.Tensor'>

Downloading (...) Pooling/confusion: 100%          190/190 [00:00<00:00, 13.0kB/s]

```

## ▼ The Model

```

import torch
import torchvision
from torch import mul, cat, tanh, relu

class VQA_v2(torch.nn.Module):
    def __init__(self, embedding_size, num_answers):
        super(VQA_v2, self).__init__()

        # The Image network which processes image and outputs a vector of shape (batch_size x 32)
        resnet = torchvision.models.resnet50(pretrained=True)
        num_fts = resnet.fc.in_features
        resnet.fc = torch.nn.Sequential(
            torch.nn.Linear(num_fts, 512),
            torch.nn.Tanh(),
            torch.nn.Linear(512, 128),
            torch.nn.Tanh(),
            torch.nn.Linear(128, 32)
        )

```

```

self.mdl = resnet

# The question network processes the question and outputs a vector of shape (batch_size x 32)
self.fc2 = torch.nn.Linear(embedding_size, 64)      # (384, 64)
self.fc3 = torch.nn.Linear(64, 32)                  # (64, 32)

# Layers for Merging operation
self.fc4 = torch.nn.Linear(64, 32)
self.fc5 = torch.nn.Linear(32, num_answers)

def forward(self, x, q):
    # The Image network
    x = self.mdl(x)                                # (batch_size, 32)

    # The question network
    act = torch.nn.Tanh()
    q = act(self.fc2(q))                            # (32, 32)
    q = act(self.fc3(q))                            # (32, 32)

    # Merge -> output
    out = cat((x, q), 1)                            # concat function
    out = act(self.fc4(out))                         # activation
    out = self.fc5(out)                             # output probability
    return out

```

## ▼ Custom Dataset

```

from torch.utils.data import Dataset

class CustomDataset(Dataset):
    def __init__(self, img, txt, ans):
        self.img = img
        self.txt = txt
        self.ans = ans

    def __len__(self):
        return len(self.ans)

    def __getitem__(self, idx):
        ans = self.ans[idx]
        img = self.img[idx]
        txt = self.txt[idx]
        return img, txt, ans

```

## ▼ Train and evaluate loops

```

def train_loop(model, optimizer, criterion, train_loader):
    model.train()
    model.to(device)
    total_loss, total = 0, 0

    for image, text, label in train_loader:
        # get the inputs; data is a list of [inputs, labels]
        image, text, label = image.to(device), text.to(device), label.to(device)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        output = model.forward(image, text)
        loss = criterion(output, label)
        loss.backward()
        optimizer.step()

        # Record metrics
        total_loss += loss.item()
        total += len(label)

    return total_loss / total

```

```
def validate_loop(model, criterion, valid_loader):
    model.eval()
    model.to(device)
    total_loss, total = 0, 0

    with torch.no_grad():
        for image, text, label in testloader:
            # get the inputs; data is a list of [inputs, labels]
            image, text, label = image.to(device), text.to(device), label.to(device)

            # Forward pass
            output = model.forward(image, text)

            # Calculate how wrong the model is
            loss = criterion(output, label)

            # Record metrics
            total_loss += loss.item()
            total += len(label)

    return total_loss / total
```

## ▼ WandB for Logging

```
!pip install WandB
import wandb
# login with your API key from wandb account
!wandb login --relogin

Requirement already satisfied: WandB in /usr/local/lib/python3.10/dist-packages (0.15.12)
Requirement already satisfied: Click!=8.0.0,>=7.1 in /usr/local/lib/python3.10/dist-packages (from WandB) (8.1.7)
Requirement already satisfied: GitPython!=3.1.29,>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from WandB) (3.1.40)
Requirement already satisfied: requests<3,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from WandB) (2.31.0)
Requirement already satisfied: psutil>=5.0.0 in /usr/local/lib/python3.10/dist-packages (from WandB) (5.9.5)
Requirement already satisfied: sentry-sdk>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from WandB) (1.32.0)
Requirement already satisfied: docker-pycreds>=0.4.0 in /usr/local/lib/python3.10/dist-packages (from WandB) (0.4.0)
Requirement already satisfied: PyYAML in /usr/local/lib/python3.10/dist-packages (from WandB) (6.0.1)
Requirement already satisfied: pathtools in /usr/local/lib/python3.10/dist-packages (from WandB) (0.1.2)
Requirement already satisfied: setproctitle in /usr/local/lib/python3.10/dist-packages (from WandB) (1.3.3)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from WandB) (67.7.2)
Requirement already satisfied: appdirs>=1.4.3 in /usr/local/lib/python3.10/dist-packages (from WandB) (1.4.4)
Requirement already satisfied: protobuf!=4.21.0,<5,>=3.19.0 in /usr/local/lib/python3.10/dist-packages (from WandB) (3.20.3)
Requirement already satisfied: six>=1.4.0 in /usr/local/lib/python3.10/dist-packages (from docker-pycreds>=0.4.0->WandB) (1.16.0)
Requirement already satisfied: gitdb<5,>=4.0.1 in /usr/local/lib/python3.10/dist-packages (from GitPython!=3.1.29,>=1.0.0->WandB) (4.0.5)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.0.0->WandB) (3.2.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.0.0->WandB) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.0.0->WandB) (2.0.4)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.0.0->WandB) (2023.7.22)
Requirement already satisfied: smmap<6,>=3.0.1 in /usr/local/lib/python3.10/dist-packages (from gitdb<5,>=4.0.1->GitPython!=3.1.29,>=1.0.0->WandB) (5.0.0)
wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: https://wandb.me/wandb-server)
wandb: You can find your API key in your browser here: https://wandb.ai/authorize
wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit:
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc

# set path
from pathlib import Path
from google.colab import drive
drive.mount('/content/drive')
project_path = '/content/drive/MyDrive/FSDL/'
project_path = Path(project_path)
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True)

## ▼ Training Run ( + Dataloaders and Hyper parameters)

```
from torch.utils.data import DataLoader

from tqdm.notebook import tqdm
# WandB - Initialize a new run
wandb.init(project='easy-vqa-pytorch',
           name='ResNet-4FC-SBERT-Tanh-20Epoch-Med',
```

```

notes='ninth run')

# WandB – Config is a variable that holds and saves hyperparameters and inputs
config = wandb.config # Initialize config
config.batch_size = 32 # input batch size for training (default: 64)
config.test_batch_size = 32 # input batch size for testing (default: 1000)
config.epochs = 40 # number of epochs to train (default: 10)
config.lr = 0.01 # learning rate (default: 0.01)
config.momentum = 0.5 # SGD momentum (default: 0.5)
config.no_cuda = False # disables CUDA training
config.log_interval = 10 # how many batches to wait before logging training status

if torch.cuda.is_available(): device = torch.device("cuda:0")
kwargs = {'num_workers': 1, 'pin_memory': True} if torch.cuda.is_available() else {}

# Now we load our training and test datasets initialize the train, validation, and test data loaders

train_dataset = CustomDataset(train_X_ims, train_X_seqs, train_Y)
test_dataset = CustomDataset(test_X_ims, test_X_seqs, test_Y)
trainloader = DataLoader(train_dataset, shuffle=True, batch_size=config.batch_size)
testloader = DataLoader(test_dataset, batch_size=config.test_batch_size)

# Initialize our model, recursively go over all modules and convert their parameters and buffers to CUDA tensors (if device is set)
model = VQA_v2(embedding_size = 384, num_answers = num_answers).to(device)
criterion = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=config.lr,
                             momentum=config.momentum )

# WandB – wandb.watch() automatically fetches all layer dimensions, gradients, model parameters and logs them automatically to wandb
# Using log="all" log histograms of parameter values in addition to gradients
wandb.watch(model, log="all")
train_losses, valid_losses = [], []

for epoch in range(config.epochs):
    train_loss = train_loop(model, optimizer, criterion, trainloader)
    valid_loss = validate_loop(model, criterion, testloader)

    tqdm.write(
        f'epoch #{epoch + 1:3d}\ttrain_loss: {train_loss:.2e}\tvalid_loss: {valid_loss:.2e}\n',
    )

    train_losses.append(train_loss)
    valid_losses.append(valid_loss)

wandb.log({
    "Epoch": epoch,
    "Training Loss": train_loss,
    "Validation Loss": valid_loss})

```

**wandb:** Currently logged in as: **yugingsu9**. Use ``wandb login --relogin`` to force r  
Tracking run with wandb version 0.15.12  
Run data is saved locally in /content/wandb/run-20231029\_215333-lab84zkc  
Syncing run **ResNet-4FC-SBERT-Tanh-20Epoch-Med** to **Weights & Biases** ([docs](#))  
View project at <https://wandb.ai/yugingsu9/easy-vqa-pytorch>  
View run at <https://wandb.ai/yugingsu9/easy-vqa-pytorch/runs/lab84zkc>  
/usr/local/lib/python3.10/dist-packages/torchvision/models/\_utils.py:208: UserWarning: warnings.warn(  
/usr/local/lib/python3.10/dist-packages/torchvision/models/\_utils.py:223: UserWarning: warnings.warn(msg)

Downloading: "<https://download.pytorch.org/models/resnet50-0676ba61.pth>" to /root/.cache/torch/hub/pytorch\_vision-0.15.1/models/resnet50-0676ba61.pth [00:01<00:00, 93.9MB/s]

epoch #	train_loss	valid_loss
1	5.28e-02	4.93e-02
2	3.86e-02	2.77e-02
3	2.36e-02	1.96e-02
4	1.68e-02	1.45e-02
5	1.33e-02	1.23e-02
6	1.17e-02	1.12e-02
7	1.08e-02	1.08e-02
8	1.03e-02	9.97e-03
9	9.94e-03	9.73e-03
10	9.66e-03	9.71e-03
11	9.48e-03	9.58e-03
12	9.36e-03	9.58e-03
13	9.22e-03	9.64e-03
14	9.07e-03	9.75e-03
15	8.96e-03	9.37e-03
16	8.84e-03	9.51e-03
17	8.75e-03	9.62e-03
18	8.65e-03	9.64e-03
19	8.60e-03	1.02e-02
20	8.47e-03	9.72e-03
21	8.36e-03	9.55e-03
22	8.30e-03	9.81e-03
23	8.18e-03	9.46e-03

## ▼ Loss charts

```
import matplotlib.pyplot as plt
%matplotlib inline
%config InlineBackend.figure_formats = ['svg']
plt.style.use('ggplot')
```

```
epoch_ticks = range(1, epoch + 2)
plt.plot(epoch_ticks, train_losses)
plt.plot(epoch_ticks, valid_losses)
plt.legend(['Train Loss', 'Valid Loss'])
plt.title('Losses')
```



```
plt.xlabel('Epoch #')
plt.ylabel('Loss')
plt.xticks(epoch_ticks)
plt.show()
```



## ▼ Validation Accuracy

```
model.eval()
model.to(device)
num_correct = 0
num_samples = 0
predictions = []
answers = []

with torch.no_grad():
    for image, text, label in testloader:
        image, text, label = image.to(device), text.to(device), label.to(device)
        probs = model.forward(image, text)

        _, prediction = probs.max(1)
        predictions.append(prediction)

        answer = torch.argmax(label, dim=1)
        answers.append(answer)

    num_correct += (prediction == answer).sum()
    num_samples += prediction.size(0)

valid_acc = (f'Got {num_correct} / {num_samples} with accuracy {float(num_correct)/float(num_samples)*100:.2f}')
print(valid_acc)

wandb.log({
    "Validation Accuracy": round(float(num_correct)/float(num_samples)*100, 2)})

Got 8566 / 9673 with accuracy 88.56
```

## ▼ Saving and Loading the model

```
torch.save(model.state_dict(), project_path/'Resnet-Sbert-40')

model = VQA_v2(embedding_size = 384, num_answers = 13)
model.load_state_dict(torch.load(project_path/'Resnet-Sbert-40')) # Im loading my best model
model.eval()
```

```

        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)
  (1): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (2): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Sequential(
  (0): Linear(in_features=2048, out_features=512, bias=True)
  (1): Tanh()
  (2): Linear(in_features=512, out_features=128, bias=True)
  (3): Tanh()
  (4): Linear(in_features=128, out_features=32, bias=True)
)
(fc2): Linear(in_features=384, out_features=64, bias=True)
(fc3): Linear(in_features=64, out_features=32, bias=True)
(fc4): Linear(in_features=64, out_features=32, bias=True)
(fc5): Linear(in_features=32, out_features=13, bias=True)
)

```

```

from urllib.request import urlopen
from PIL import Image

```

```

def load_and_process_image_url(url):
    # Loads image from path and converts to Tensor, you can also reshape the im
    im = Image.open(urlopen(url))
    im = F.to_tensor(im)
    return im

```

```

url = "https://www.nicepng.com/png/detail/16-163438_circle-clipart-sky-blue-clip-art-blue-circle.png"
image = load_and_process_image_url(url)
image = image.unsqueeze(0)

```

```

text = 'What shape is this?'
text = st_model.encode(text)
text = torch.tensor(text, dtype=torch.float)
text = text.unsqueeze(0)

```

```

probs = model.forward(image, text)
answer_idx = torch.argmax(probs, dim=1) # get index of answer with highest probability
answer_text = [all_answers[idx] for idx in answer_idx] # convert index to answer text
print(answer_text)

```

```
['triangle']
```